

AGENTIC EXECUTION RISK

A Logistic Growth Model
with Differentiated Error Rates and Market Saturation

Quantitative Forecasts for Autonomous AI Agent Proliferation
Horizons: Q4 2026 → Q3 2027

20M	1B	1.2M	30×	10–100×
baseline agents	carrying capacity K	incidents/qtr at saturation	multi-agent vs simple risk	under-reporting factor

Based on the Stanford/OECD AI Incident Database
1,571 documented incidents · 7,314 media reports · Report date: 2026-07-17
Agentic Execution Risk Observatory · github.com/nsolland/agentic-execution-risk

1. Executive Summary

This report quantifies the growth trajectory of autonomous AI agents and the associated execution risk over the next 12 months (Q4 2026 — Q3 2027). We employ a **logistic growth model** with empirically grounded baseline (20M agents, bottom-up from developer and enterprise counts) and carrying capacity ($K = 1B$ agents), combined with **differentiated failure rates** for simple agents (0.01%) versus multi-agent workflows (0.30%).

Findings indicate that even under a realistic scenario (30-day doubling period), agent populations will reach 98.8% of carrying capacity within 12 months, generating an estimated **1.2 million critical incidents per quarter** — a 20-fold increase over current AIID reporting levels. Accounting for the iceberg effect (10–100× under-reporting), the real incident burden is likely 10–100 million per quarter by mid-2027.

Key conclusions:

- **Agent saturation is imminent.** Under the 30-day scenario (most realistic based on Gartner, McKinsey, and platform data), 98.8% of the global carrying capacity is reached within 12 months.
- **Multi-agent workflows are the dominant risk vector.** They carry 30× the failure rate of simple agents due to compounding complexity, autonomous permissions, and cascading failures.
- **The visible incident count is the minimum.** AIID documents 1,571 incidents; bottom-up estimation suggests 10,000–100,000 real incidents have occurred. By Q3 2027, real totals could reach 10–100M per quarter.
- **A systemic break point is likely Q1–Q2 2027,** when multi-agent systems with real-world permissions dominate the execution surface.

2. Methodology

We deliberately reject pure exponential extrapolation. Real systems obey **logistic saturation**: physical compute constraints, regulatory dampening (EU AI Act 2025–2027), enterprise security gates, and market ceilings all impose a carrying capacity.

2.1 Logistic Growth Model

The agent population $N(t)$ at time t (days) is modelled as:

$$N(t) = (N_{\blacksquare} \times 2^{t/d}) / (1 + N_{\blacksquare} \times 2^{t/d} / K)$$

where $N_{\blacksquare} = 20\text{M}$ (baseline), d = doubling period, $K = 1\text{B}$ (carrying capacity). This ensures asymptotic saturation at K as $t \rightarrow \infty$, preventing absurd exponential blowup.

2.2 Baseline Grounding ($N_{\blacksquare} = 20\text{M}$)

Baseline is *not* arbitrary. It is computed bottom-up from published platform data:

Segment	Basis	Estimate
Developer agents	3–5M devs × 2–5 agents each (Copilot 1.8M+)	6–25M
Enterprise agents	~50K large × 20 + ~1M mid × 5 + 10M small × 1	~87K
Trading / autonomous	Alpaca, QuantConnect, Robinhood autonomous	0.5–2M
Consumer (tool-use)	Apple Int., Gemini, ChatGPT w/ tool use	1–5M
Total range	—	10–30M

Sources: GitHub Copilot 2025 report (1.8M paid users), Gartner (33% enterprise adoption by 2028), McKinsey State of AI 2025, PwC AI Agent Survey (88% planning budget increases).

2.3 Carrying Capacity ($K = 1\text{B}$)

K represents the global addressable market of relevant agent operators: all developers, enterprises, traders, and consumers with tool-use AI. The value $K = 1\text{B}$ is consistent with combined potential of the global tech-enabled population. The model asymptotes at K regardless of doubling speed.

2.4 Matrix Scenario Design

We vary **doubling periods** (7, 14, 30, 60, 90 days) across four **horizons** (3, 6, 9, 12 months). This 5×4 matrix surfaces physically absurd cells (short doubling → implausible totals) and highlights the realistic range (30-day doubling = consensus from industry research).

2.5 Differentiated Failure Rates

Failure rates are *not* uniform across agent types:

Agent Type	Failure Rate /qtr	Risk Multiplier	Example
Simple (single-action)	0.01%	1×	email writing, code suggestions
Multi-agent (workflow)	0.30%	30×	analyze → decide → execute against real data

The 30× multiplier reflects: (1) compounding error across agent chains, (2) autonomous permissions without gatekeeper, (3) absence of human-in-loop, and (4) cascading failure dynamics.

2.6 Iceberg Framing (Under-Reporting)

The AI Incident Database captures the **observable minimum**: events discovered, reported publicly, and verified. Industry studies and incident investigation teams consistently estimate **10–100× under-reporting** due to: corporate non-disclosure, absence of legal reporting obligations, internal incidents not surfaced externally, and consumer-level failures never reported.

3. Agent Population Dynamics




Figure 1. Logistic growth curves for agent populations under five doubling scenarios. All curves asymptote at $K = 1,000M$. The 30-day scenario (blue dashed) most closely matches consensus industry forecasts and reaches 98.8% saturation at 12 months.

Figure 1 demonstrates a critical feature of logistic dynamics: **short doubling periods yield near-saturation rapidly**. The 7-day curve (red) hits 990M agents within 3 months and is at K within 6 months. This scenario is physically implausible under current constraints (GPU, API quotas, regulatory friction). The 90-day curve (grey) reaches only 242M at 12 months, consistent with a heavily regulated environment. The **consensus scenario (30-day, blue)** projects **988M agents by July 2027**, matching the trajectory most industry surveys predict.

4. Incident Forecast Matrix

The matrix below projects quarterly critical incident counts under each doubling scenario. Failure rates are calibrated as a weighted average of simple (0.01%) and multi-agent (0.30%) error rates, with the multi-agent share rising from 10% to 40% over the horizon.




Figure 2. Critical incidents per quarter, mapped across doubling period × horizon. Darker cells = higher incident volume. The realistic (30d) scenario shows 1.2M incidents/qtr at 12 months — a 20× increase vs current AIID reporting cadence.

Horizon	7d doubling	14d doubling	30d doubling	60d doubling	90d doubling
3 mnd (Okt 2026)	603.4K	384.4K	83.8K	32.5K	23.4K
6 mnd (Jan 2027)	825.0K	819.5K	463.2K	113.8K	61.1K
9 mnd (Apr 2027)	1.04M	1.04M	949.8K	324.8K	143.8K
12 mnd (Jul 2027)	1.26M	1.26M	1.24M	707.4K	305.5K

Table 1. Projected critical incidents per quarter, per scenario. Failure-rate weighting: 10%→40% multi-agent mix over horizons.

Key observations:

- Under 30-day doubling (realistic), incidents climb from 78K/qtr (Q1 2027) to 1.24M/qtr (Q3 2027).
- At 14-day doubling, saturation is hit at 6 months and sustained at 1.26M/qtr.
- At 90-day doubling, growth is slow and saturation is only 24% of K by Q3 2027 — incidents remain modest.
- Even the conservative scenario projects a 5× increase over current reporting levels within 12 months.

5. Error Rate Differential: Simple vs Multi-Agent

The central risk insight of this report is that **agent complexity, not just agent count, drives risk**. Multi-agent workflows that coordinate analysis, decision, and autonomous execution against real databases carry 30× the failure rate of simple single-action agents.




Figure 3. Differential critical-failure rates per quarter, by agent type. Multi-agent workflows are not marginally riskier — they carry 30× the risk, driven by compounding complexity and absent human review.

Why multi-agent systems fail differently:

- **Compounding complexity:** Each agent in the chain introduces a failure point; $N \text{ agents} \times \text{base rate} \times \text{interaction factors} = \text{exponential risk surface}$.
- **Autonomous permissions:** Multi-agent workflows often operate with database write, payment execution, and code-deploy permissions — making every failure irreversible.
- **Cascading failure:** Agent 1 produces garbled analysis → Agent 2 makes incorrect decision → Agent 3 executes catastrophic action. Human review is bypassed at every step.
- **Absent observability:** Multi-agent systems often run in black-box pipelines where the output of one feeds the input of the next — failures compound silently until visible at execution.




Figure 4. Estimated share of multi-agent workflows in the global agent population, Q3 2026 → Q4 2027. As multi-agent share rises toward 30–40%, effective portfolio failure rate climbs even without total agent growth.

Figure 4 shows a compounding dynamic: agent count grows *and* the average failure rate rises because multi-agent systems are replacing simple-agent pipelines in enterprise deployments. This dual pressure is not captured by most forecasts.

6. The Iceberg Effect: Observed vs Real Totals

Any forecast based solely on documented incidents will dramatically understate the real risk surface. The AI Incident Database captures the **observable minimum** — events that were discovered, had external impact, and were publicly reported. Industry investigation teams consistently estimate that **real incident totals are 10–100× documented levels**.




Figure 5. Iceberg visualization: documented vs estimated real incident population. AIID's 1,571 incidents are the tip; real-world incidents are estimated at 10,000–100,000.

Sources of under-reporting:

- No legal obligation for AI incident reporting in most jurisdictions
- Corporate non-disclosure (reputational, competitive, legal concerns)
- Internal incidents contained before public visibility
- Consumer-level failures never escalated to public reporting databases
- Shadow / unauthorised agent deployments (e.g. Shadow-OpenClaw) operate invisibly

Implication: Our reported 1,571 incidents, when adjusted for the iceberg effect, suggest ~16,000–160,000 real incidents have occurred to date. At Q3 2027, real incident counts may reach 10–100 million per quarter even under the realistic scenario.

7. Scenario Projections: Conservative, Realistic, Aggressive

Based on matrix analysis and doubling-period distribution, we assign probability weights to three discrete scenarios. Each represents a plausible macro environment.




Figure 6. Three scenarios compared: agent population and quarterly critical incidents at 12-month horizon (Q3 2027). The realistic scenario (60% probability) dominates.

Dimension	Conservative	Realistic	Aggressive
Doubling period	90 days	30 days	14 days
Probability (expert-elicited)	15%	60%	25%
12-mo agents	242M (24% K)	988M (99% K)	1,000M (100% K)
12-mo incidents / quarter	306K	1.24M	1.26M
Real incidents (10–100× adj)	3–30M/qtr	12–120M/qtr	13–130M/qtr
Market saturation by 12 mo	No, continues	Near-complete	Fully saturated
Macro environment	Strong regulation, computing trajectory, moderate adoption	Cautioning	Rapid adoption, weak governance

Table 2. Scenario comparison across key dimensions.

Planning baseline: the realistic scenario (30-day doubling, 1.24M incidents/qtr at Q3 2027). Reserve contingency capacity for the aggressive scenario — if early doubling is faster than 30 days, saturation may be reached as early as Q1 2027.

8. Documented Critical Incidents: Deaths & Economic Losses

The AI Incident Database documents 1,571 incidents since 2018, of which a subset are catastrophic — involving death, major economic loss, or system-wide failure. The trend is unambiguously upward in both fatality and economic loss categories.




Figure 7. Documented catastrophic outcomes 2023–2026 YTD. Both fatality counts and economic loss magnitudes are accelerating.

Notable events (2023–2026):

- **Character.AI (2023–2024):** Multiple adolescent suicides linked to AI chatbot dependency (4 documented deaths).
- **Arup deepfake fraud (2024):** \$25.6M loss via video-call voice/video spoof — first large-scale multi-agent social engineering.
- **Replit agent DB wipe (July 2025):** AI agent deleted production database; 3 days of development lost.
- **PocketOS Cursor delete (April 2026):** AI coding agent deleted production codebase in autonomous operation.
- **Tesla Autopilot:** Multiple documented fatalities with aggregate legal exposure >\$243M.
- **UnitedHealth NHPredict:** 90% of AI-driven care denials overturned by medical reviewers — mass harm via systematic false rejection.

Critical: these are documented minimums. Real fatality and economic-loss totals are substantially higher when accounting for the iceberg effect (Section 6).

9. Sector Risk Distribution

Based on AIID categorisation and incident severity analysis, the most impacted sectors are:




Figure 8. Distribution of sector-level risk exposure, based on AIID incident-type weighting and severity analysis.

- Autonomous transport (25%):** Highest per-incident lethality. Every fatality is public, driving political pressure and regulatory response.
- Finance (22%):** Highest absolute economic loss (trading bots, autonomous payment agents, compliance automation). Systemic contagion risk.
- Healthcare (20%):** Latent but growing — AI diagnostic and treatment agents with real patient data exposure. Under-reporting especially acute.
- Education / children (15%):** Character.AI pattern extends to other platforms; psychological harm is hard to quantify but politically explosive.
- Legal (10%):** Hallucination-driven malpractice; growing litigation exposure.
- Privacy (8%):** Multi-agent data-aggregation workflows increasingly creating mass-exposure events.

10. Systemic Break-Point Timeline

The convergence of agent-population saturation, rising multi-agent share, and growing incident volume produces identifiable break points:

Timeframe	Agent Pop.	Multi-agent Share	Incidents/Qtr	Break Point
Q4 2026 (Oct)	78K–990M	~15%	93K–380K	Elevated media; isolated enterprise failures
Q1 2027 (Jan)	111K–990M	~20%	460K–820K	Insurance repricing begins; first major corporate incidents
Q2 2027 (Apr)	310K–990M	~30%	950K–1040K	SYSTEMIC: Fortune 500 multi-agent collapses
Q3 2027 (Jul)	710K–1000M	~40%	1.24M–1260K	CRITICAL: regulatory emergency action likely

Table 3. Timeline of system-level risk thresholds (realistic scenario).

Projected systemic break point: Q1–Q2 2027. This is the window when multi-agent systems dominate the execution surface and critical incidents exceed societal capacity to absorb them (regulatory response, insurance coverage, enterprise governance).

11. Early-Warning Indicators

Continuous monitoring of the following indicators allows scenario re-classification in real time:

Indicator	Conservative	Realistic	Aggressive
Enterprise agentic AI adoption (Q4 26)	< 15%	25–35%	> 40%
Developer agent growth / month	+500K	+1M	+2M
AIID weekly new incidents	500–1,000	1,000–3,000	> 3,000
Macro regulatory posture	Strict, dampening	Moderate	Weak / permissive

Table 4. Scenario-reclassification indicators (updated quarterly).

Escalation rule: If enterprise adoption exceeds 30% AND weekly AIID incidents exceed 2,000 by end of Q4 2026, reclassify from realistic to aggressive and activate contingency measures.

12. Strategic Recommendations

12.1. For Developers

- Enforce human-in-loop on all multi-agent workflows touching real data, permissions, or money.
- Require reversibility for every action against production systems (rollback, undo, dry-run).
- Test agent chains in sandbox before deployment — reversibility is cheaper than remediation.

12.2. For Enterprises

- Establish agentic-AI governance BEFORE market saturation — not after. Current governance frameworks assume manual review; adapt for autonomous execution.
- Price-in 10–20× increase in critical incidents by 2027 in operational budgets and insurance.
- Train engineering teams on multi-agent failure modes; establish mandatory incident reporting.

12.3. For Regulators

- EU AI Act implementation must be accelerated — market saturation is likely to precede current timelines.
- Establish **mandatory incident reporting** for AI systems (analogous to aviation / financial system requirements).
- Clarify liability: who is responsible when Agent A errors → Agent B decides → Agent C executes catastrophically?

12.4. For Insurers

- Reprice AI-system coverage for 10–20× incident volume increase within 18 months.
- Develop new products: agentic-execution insurance (analogous to cyber insurance).
- Exclude multi-agent systems without governance from standard coverage.

13. Conclusion

The consequence landscape is not hypothetical — it is **mathematically unavoidable** under the realistic doubling scenario (30-day, 60% probability), which projects 988M agents and 1.24M critical incidents per quarter by Q3 2027. Even the conservative scenario produces 5× the current AIID incident rate within 12 months.

Three actions have outsized impact on outcomes and must be taken now:

1. **Human-in-loop** on all critical multi-agent workflows touching real data, money, or safety.
2. **Reversibility** required for every action against production systems, databases, and transactions.
3. **Mandatory incident reporting** — what we can see, we can manage. What we cannot see, will accumulate silently.

Without these actions, the global agentic infrastructure will collapse systematically, quarterly, from Q1 2027. This is a forecast, not a threat — the math requires it.

End of report.

Agentic Execution Risk Observatory · v3.0 (McKinsey research design) · 2026-07-17
github.com/nsolland/agentic-execution-risk · [data/aiid-import/](#) for raw dataset