

May 2026

THE
FUTURE
SOCIETY

The Case for Cross-Border AI Incident Infrastructure

Caio Vieira Machado, George Gor
and Omer Bilgin

Contact: info@thefuturesociety.org

Cite as:

Caio Vieira Machado, George Gor, and Omer Bilgin, “The Case for Cross-Border AI Incident Infrastructure” (The Future Society 2026).

© 2026 by The Future Society. Cover and back image created with Midjourney.

This work is licensed under a [Creative Commons Attribution-NonCommercial 4.0 International License](https://creativecommons.org/licenses/by-nc/4.0/).



Table of Contents

Table of Contents	3
Executive Summary	5
1. Introduction	7
2. The Need for Cross-Border AI Incident Infrastructure	10
2.1. Harm is Rising Across Severe Risk Categories	10
2.2. The Systems Producing These Harms Span Jurisdictions	14
2.3. Detection and Response for AI incidents Are Currently Reactive by Design	15
3. The Architecture Required and How It Can Be Secured	17
4. Conclusion	30
Acknowledgments	32
References	33

Executive Summary

General-purpose artificial intelligence (AI) systems are producing harm at scale, on a development timeline that has consistently outrun the institutions meant to govern them. AI systems influence users' beliefs and behavior in ways that cause severe harm, for example by reinforcing suicidal ideation through extended conversation without triggering any external safeguard. AI systems have also been used to generate hundreds of thousands of instances of child sexual abuse material, publicly spread across multiple jurisdictions through social media platforms, and their interaction designs have been worsening mental health among both youth and adult populations. Recently, some of the most advanced and capable ("frontier") AI systems have exhibited capabilities that can be used maliciously to plan and execute sophisticated cyber operations against critical infrastructure in multiple jurisdictions at once.

This report makes the case for **cross-border AI incident infrastructure: the institutional, technical, and legal arrangements that would allow national and international authorities to detect, share, and act on AI incidents in coordination**. Three conditions make this work urgent.

- **Harm is rising.** Public databases that monitor AI incidents and hazards, such as the one run by the Organization for Economic Cooperation and Development (OECD), show a significant increase in reported incidents over time, particularly those relating to cybercrime and flare-ups in mental health issues in users.
- **The systems producing these harms cross borders by design.** A model built in one country reaches users in dozens of others within days. Developers assemble training data in one jurisdiction, train models in another, and deploy worldwide at once. However, regulators, even in the countries where the largest AI developers are based, hold only partial capacity to investigate, let alone respond.
- **Detection and response are reactive by design.** Existing incident monitoring mechanisms rely on voluntary disclosure by frontier AI companies, journalistic coverage, and public outcry, mechanisms that end up only surfacing patterns after serious harm has already occurred. This makes anticipatory governance extremely difficult.

The six architectural components required for this infrastructure, and concrete actions that can be taken to secure them, include the following:

- 1) **Internationally interoperable protocols for incident documentation and reporting.** Governments and industry actors should adopt incident documentation structures and protocols that are universally compatible with one another, so that incident data collected in one jurisdiction becomes evidence that other jurisdictions can use.
- 2) **Mandatory incident documentation, disclosure, and preparedness obligations on frontier AI companies.** Governments should impose statutory incident reporting obligations on frontier AI companies and require them to publish policies for preventing, preparing for, and responding to serious AI incidents.

- 3) **Domestic incident data access and institutional coordination.** Governments should grant national AI Safety Institutes, or functionally equivalent bodies, statutory access to incident data held by AI deployers and sectoral regulators, paired with formal information-sharing arrangements between such bodies and regulatory authorities.
- 4) **Adequate international incident analysis capacity and expertise.** International standards development organizations should develop a standardized methodology for AI incident root-cause and supply-chain analysis, with multilateral institutions such as the OECD and the United Nations (UN) advocating for its adoption and national bodies investing in the technical capacities needed to apply it.
- 5) **Cross-border incident-sharing channels and joint investigation mechanisms.** National AI governance bodies, cybersecurity agencies, and sector-specific regulators should be empowered to share incident intelligence with counterparts abroad, for example through the International Network for Advanced AI Measurement, Evaluation and Science, or through a new international mechanism that could be proposed via the UN's Global Dialogue on AI Governance.
- 6) **Internationally distributed incident detection, analysis, and response capacity.** Multilateral institutions should help extend incident infrastructure to jurisdictions who have inadequate domestic incident monitoring and management capacities through the establishment of bilateral or multilateral intelligence-sharing, capacity-building, and incident response coordination partnerships between countries that already possess relatively mature incident infrastructure and those that do not.

The Policy Window of Opportunity

Jurisdictions around the world have begun requiring frontier AI developers to report serious safety and security incidents to designated authorities and publish how they manage severe risks. So far, the European Union's AI Act, emerging state legislation in the United States such as California's Transparency in Frontier Artificial Intelligence Act and New York's RAISE Act, and South Korea's AI Basic Act now require incident reporting on tight timelines, and most also require large developers to publish the frameworks governing how they assess and mitigate severe risks. While California and South Korea's regime took effect in early 2026, the EU's serious incident reporting obligations apply from August 2026, and New York's RAISE Act takes effect in January 2027.

With sustained effort, countries should move from tracking AI incidents inside their own borders to seeing what the companies producing these systems already see, and responding to what they find together. The work is more achievable now, at a time when domestic regimes are still being developed and enforced.

1. Introduction

On February 10, 2026, a devastating mass shooting in Tumbler Ridge, British Columbia, claimed the lives of eight people, including five young children, and left 27 others injured ([Judd 2026](#)). Months before the attack, the perpetrator's ChatGPT account had been banned by the US-based developer, OpenAI, due to troubling chat logs featuring scenarios of gun violence. The perpetrator had expressed their intent to kill during their interactions with ChatGPT, and had even created a new account after the initial ban to use the chatbot to plan their attack ([Yousif 2026](#)). Approximately a dozen US-based staff members debated alerting Canadian authorities, actively identifying the user's behavior as an indication of potential real-world violence. However, **operating without the necessary cross-border infrastructure or standardized escalation protocols, company leadership ultimately decided not to contact law enforcement**, determining internally that the activity did not meet their threshold for a credible or imminent plan for serious physical harm ([Cress 2026](#)). It was only after the casualties occurred that Canada's AI Minister Evan Solomon summoned OpenAI executives to Ottawa to discuss safety protocols and thresholds for police escalation ([Yousif 2026](#)).

This tragedy was facilitated by an AI system and was detected early on, internally, by its deployer. As such, this disaster represents a profound, preventable cross-border intelligence failure, one that highlights the urgent need for coordinated AI incident reporting and response. As this case vividly illustrates, the evidence required to detect and effectively respond to severe AI-enabled harm often sits inside companies that develop the most capable (i.e. “frontier”) general-purpose AI (GPAI) systems, while the institutions empowered to act on that evidence are forced to rely on whatever portion these companies voluntarily choose to disclose.

Events like that of Tumbler Ridge reflect the growing scale and intensity of harms that AI development and use can cause in societies around the world, and how the risks and effects are interconnected across multiple geographies. **These are examples of what has been called severe AI incidents**, which have been spreading around the world. There are ongoing definitional disputes around what exactly “AI incidents” and “AI hazards” entail. In such discussions, there are certain terms that are often brought up, including *accidents*, *crises*, *disasters*, *catastrophes*, *emergencies*, *near-misses*, *flaws*, *weaknesses* and *vulnerabilities* ([OECD 2024](#)).

Some experts are currently discussing the appropriateness of these terms and how they relate to AI incidents and hazards. While we acknowledge that these disputes exist and that other literature might distinguish these terms, for the purposes of this report, we group them all within the categories of “AI incidents” and “AI hazards”. We use **AI incidents** to refer to events or circumstances where the development, use, or malfunction of AI systems directly or indirectly leads to real-world harm to individuals, organizations, communities, property, or the environment ([OECD 2024](#); [EU AI Act Article 3\(49\)](#)). Relatedly, we use **AI hazards** to refer to events, circumstances, built-in system features, and other relevant findings that suggest that the

development, use, or malfunction of an AI system could plausibly lead to incidents, such as near-misses or the emergence of novel capabilities, signaling potential risks that have not yet materialized into real-world harm. If left unaddressed, AI hazards and incidents can rapidly escalate into national or international emergencies, disasters, and crises ([Gor & Iliadis 2025](#)).

Looking at the narrower space of severe risks and harms, these cases reflect emerging patterns emanating from GPAI development and deployment, where documented harms span multiple risk categories, from harmful manipulation and child safety to autonomous cyber operations, loss-of-control behaviors, and CBRN uplift ([Ileri 2025](#); [OECD 2026](#); [UK AISI 2026](#); [Shane et al. 2026](#); [Bengio et al. 2026](#)). These harms cross borders by construction. AI systems today are built in a small number of jurisdictions, distributed worldwide through application programming interfaces and downstream products, and integrated across consumer applications and critical infrastructure within months of release. Governance, meanwhile, sits within national jurisdictions, with authority over the firms operating inside their borders and with reach into events elsewhere that depends on the cooperation of other governments and the developers themselves. **This mismatch between the geography of AI development and the geography of AI governance defines the challenge that this report addresses.**

National AI Safety Institutes (AISIs) and functionally equivalent institutions have built the technical capacity to evaluate the most capable AI systems and have begun to coordinate ([NIST, 2026](#)). The successive AI Summits have generated diplomatic consensus that AI risks require sustained international attention ([UK Government 2023](#); [Republic of Korea 2024](#); [French Republic 2025](#); [Government of India 2026](#)). The United Nations (UN) has operationalized the Independent International Scientific Panel on AI ([UN 2025a](#)) and the Global Dialogue on AI Governance ([UN 2025b](#)).

Each of these efforts addresses a specific component of the cross-border governance problem within a specific institutional space. For example, these institutions address challenges through capacity building for technical safety and security, or establish diplomatic baselines that continue to inform domestic legislation in several jurisdictions. The global AI governance system should build on existing insufficiencies over time, aiming at establishing guardrails to guarantee public safety and wellbeing, as well as international security.

To that end, multilateral organizations such as the Organization for Economic Co-operation and Development (OECD) maintain the AI Incidents and Hazards Monitor (OECD.AIM), which catalogues documented harms across jurisdictions and recorded 598 incidents in January 2026 alone ([OECD 2026](#); [Jeanmaire and Boger 2025](#)). Its Common Reporting Framework (CRF) offers a foundation for harmonized national reporting ([OECD 2025b](#)). The G7 Hiroshima AI Process sets out principles for developers of the most capable AI systems, covering risk management, vulnerability handling, transparency, and incident reporting, with implementation monitored by the OECD ([OECD 2025a](#)).

However, recent incidents are signals for countries to collectively and urgently establish stronger, better coordinated means of mitigating and preventing future cases. In this report, we make the case that effective governance of AI systems requires **cross-border AI incident infrastructure: the institutional, technical, and legal arrangements through which AI incidents and hazards can be detected, analyzed, shared, and acted upon across the jurisdictions where those AI systems are deployed**. This architecture would be poised to detect and respond to a variety of AI events, ranging from existing risks all the way to full-blown AI disasters. As mentioned already, for the purposes of our work, we have grouped these into the existing terminology of incidents and hazards proposed by the OECD ([OECD 2025b](#)). Such infrastructure is essential both for responding to incidents when they occur and for preventing localized failures from cascading across systems, sectors, and borders.

Over the past year, the authors of this report and The Future Society have worked on the institutional and analytical foundations of cross-border incident infrastructure. That work has involved a Main Summit session on [interoperable AI incident reporting and monitoring](#) at the 2026 India AI Impact Summit, a closed-door dialogue on [incident prevention and preparedness](#) at the Seventh Athens Roundtable, case studies for the GPAI/OECD.AI Expert Workshop on AI Incidents, and jurisdiction-specific recommendations for the [European Union](#), [United States](#), and [Canada](#). These recommendations are aimed at establishing policies that would help avert serious AI incidents, including policies for accidents, emergencies, or crises, and incident and emergency response frameworks.¹ The conclusion of that work is that **international AI governance is missing arrangements for detecting, sharing, and acting on AI incidents in a coordinated manner**.

This report sets out the architecture the field needs and the policy actions required to build and sustain such arrangements. To produce these recommendations, we benefited strongly from parallel institutional discussions on topics that lie outside the scope of this report, for example concerning the need for robust whistleblower protections at frontier AI companies, clear liability allocation for AI-related harms, and the continued development of domestic transparency regimes to produce claims worth verifying in the first place. These interventions are important and should be pursued in parallel with the infrastructure proposed here. For this report, we address decision-makers within national governments and international institutions. Frontier AI developers, civil society organizations, and academic researchers are also essential to the ecosystem this report describes. The recommendations we put forward assume their continued engagement but the actions proposed are addressed to those with authority to take them.

In what follows, [Section 2](#) establishes the empirical case for cross-border AI incident infrastructure on the grounds that harm is rising across severe risk categories and the AI systems producing it are structurally transnational. However, the current governance arrangements remain largely reactive rather than anticipatory. [Section 3](#) sets out the key architectural components of

¹ For more of this work, see [Gardhouse 2026](#); [Jeanmaire and Boger 2025](#); [Kraeva, Oueslati and Lorente 2026](#); [Gor and Iliadis 2025](#).

cross-border AI incident infrastructure, each paired with a specific actionable recommendation aimed at guiding existing institutional authorities on how to secure them in practice.

2. The Need for Cross-Border AI Incident Infrastructure

This section sets out the case for cross-border AI incident infrastructure in three parts. **Section 2.1** documents harm rising across severe risk categories, including harmful manipulation, AI-enabled cyber offense, CBRN uplift and loss of control. **Section 2.2** shows that the systems producing this harm span jurisdictions by construction, through structural features of the AI supply chain that place them beyond the reach of any single jurisdiction acting alone. **Section 2.3** argues that the governance institutions tasked with managing these systems are falling behind the capability frontier, and that the detection and response mechanisms currently available surface patterns only after serious incidents have materialized.

2.1. Harm is Rising Across Severe Risk Categories

The documented record of AI-related harm shows a substantial upward trend, consistent across the major tracking systems that measure it (see [Figure 1](#)). Each AI system is developed using different training methodologies and source datasets, and their convergence on the same finding makes the trend robust to the limits of any single approach. The OECD.AIM recorded 598 distinct incidents in January 2026 alone, compared with 299 in January 2025 and 31 in February 2020, a 19x increase in monthly reported incidents over six years ([Elagina, 2026](#); [OECD 2026](#)). The AI Incident Database (AIID) shows a comparable growth curve ([McGregor 2021](#); [AIID 2026](#)). Similarly, in the realm of AI hazards, AI-related Common Vulnerabilities and Exposures (CVEs)—publicly identified security flaws within AI systems—have surged over the past few years, with over 6,000 unique vulnerabilities identified since 2018 and 2,130 cases disclosed in 2025 alone, representing a 34.6% year-over-year increase ([Girnius et al. 2026](#)). This signals that the attack surface for cybercriminals to cause incidents concerning AI-driven fraud, destructive malware, unauthorized data access and exfiltration, and the exploitation of deployed AI agents (to strategically influence their behaviors) is rapidly expanding.

Panel A. Evolution of media-reported AI incidents and hazards over time, absolute number




Figure 1: Graph from OECD’s Trends in AI Incidents and Hazards Reported by the Media. X-axis: number of incidents and hazards. Y-axis: time period between January 2022 and September 2025 measured in months. (Source: [OECD](#))

These figures are primarily derived from media reporting and voluntary disclosure by AI developers, researchers, and users on social media platforms or stakeholder-specific reporting channels, and thus likely understate the true caseload. The OECD.AIM extracts its data from media reports, but media coverage captures only what journalists notice and consider newsworthy. There is also evidence that evaluating AI models in non-English languages reveals risks that would remain otherwise unknown ([Ileri et al. 2026](#)), yet such risks are often underexplored and discounted given frontier AI companies’ incentives to cut corners. Moreover, aggregate, society-impacting harms that unfold gradually rather than in single-event form and that are widespread in subtle ways across large groups, such as AI chatbots’ influence on mental health and misinformation discourse at scale, defy straightforward detection and reporting and are therefore systematically underrecorded ([Paeth et al. 2025](#)). After all, these harms may not involve tangible, observable harms to individuals. However, they represent emergent effects that are distributed across thousands of people across society.

Even with these caveats, across these systems and across the period they cover, these findings show that **(a)** AI-related harm has been rising substantially and that **(b)** the consistently documented incident landscape already spans several systemic risk categories, which we discuss below. Many of the documented harms are concerning, and there are early warning signals of far more serious AI incidents. Technical and policy proposals aimed at mitigating such harms are being actively developed by think tanks, proposals which we urge policymakers to pay attention to ([Kraprayoon 2024](#)). Lawmakers, in jurisdictions such as the EU and United States

(California), have also begun to respond to these concerns by setting serious-incident reporting timelines and mandating time-bound reporting for AI-related deaths or serious injuries.

Nonetheless, the report reveals important indicators over severe risks, notably:

Harmful manipulation. AI systems have influenced users' beliefs, decisions, and behaviors in ways that produced severe and sometimes irreversible harm. The signals that could detect such harm are routinely captured inside the companies producing the systems, while the channels that could route them to external authorities are still being developed. In *Raine v. OpenAI*, currently in litigation, the company's own interaction logs reportedly recorded multiple references to suicide in a single user's history that remained within the company's monitoring systems ([OpenAI 2025](#)). One of the core problems that caused the *Raine v. OpenAI* case is the persistent feature of sycophancy, which is the tendency of engagement-optimized models to validate users' existing beliefs ([Sharma et al. 2023](#)). Companies have begun responding with internal measures aimed at user wellbeing (for instance, [Anthropic 2025a](#)). However, the capacity to detect this category of harm largely sits inside developer companies, while the architecture to route those signals outward is still being built: for example under the EU Code of Practice, which now requires signatory developers to report serious incidents involving harmful manipulation, among other systemic risks, to the EU AI Office.

AI-enabled cyber offense. AI systems are now being directed to conduct sophisticated autonomous cyber operations at scale, across multiple jurisdictions at once. The category matters because the capability is operating in the wild while the response infrastructure connecting regulators across borders is still being built. Anthropic's November 2025 disclosure of the GTG-1002 operation provides an example where Claude Code had been directed by threat actors to operate as an autonomous penetration-testing agent at 80 to 90 percent autonomy across approximately thirty entities in multiple jurisdictions and critical-infrastructure sectors at once ([Anthropic 2025b](#)). Two features of the case are worth noting. First, the disclosure followed Anthropic's organizational decision, made independently of any regulatory requirement; notification to affected entities ran through the company's bilateral channels, which is the operating practice for this kind of cross-border AI incident at present. Second, the case illustrates how AI systems directed at cyber operations can outpace defenders' detection cycles, widening the offense-defense gap. Within this category, AI capability is materially outpacing both the cross-border response infrastructure built to coordinate it and the defensive cycles built to contain it.

Chemical, biological, radiological, nuclear capabilities (CBRN) uplift. AI systems have been documented to provide meaningful CBRN uplift to users, particularly non-experts, for whom AI now provides knowledge that previously required specialized training ([Luckey et al. 2025](#)). This uplift expands the population able to acquire CBRN-relevant knowledge and capability, which raises the risk profile of the category beyond what existing biosecurity instruments were designed to address. Pre-deployment CBRN evaluations are largely developer-run, with the UK AI Security Institute being the most substantively engaged AI Safety Institute outside the developers

themselves. External verification of AI safety and security claims is technically and institutionally difficult, and evaluation standards across the field are still developing ([Brundage et al. 2026](#)). The research ecosystems producing the evidence base sit in developer nations, and most other jurisdictions depend on whatever portion of those evaluations the developers choose to share ([Ileri et al. 2026](#)). Dedicated biosecurity instruments, such as synthetic-nucleotide screening and know-your-customer regimes for biological suppliers, are emerging through initiatives such as the International Biosecurity and Biosafety Initiative for Science, while verification specific to AI-enabled CBRN risk remains self-reported by developers. AI is therefore materially expanding CBRN risk, while the institutions positioned to evaluate and respond to that expansion remain a small handful concentrated in developer countries.

Loss of control. Loss-of-control behaviors that researchers documented in controlled evaluations have now appeared in deployed AI systems. The institutional capacity to detect and act on these behaviors is still being built, while the underlying capabilities continue to advance in deployed systems. In controlled evaluations, AI models in agentic configurations have attempted to strategically deceive their deployers, disable oversight mechanisms, and exfiltrate their own weights when operating in sandboxed environments. The models exhibiting these behaviors include various Claude models—including Opus 4.5 and 4.6, Sonnet 4.5 and 4.6, and Claude Code ([Anthropic 2026a](#))—Grok 4.1, Gemini 3 Pro and 3.1 Pro, GPT-5.2-Codex, GPT-5.3-Codex, and GPT-5.4 ([Shane et al. 2026](#)), with more capable models demonstrating such behaviors at higher rates and greater sophistication ([Meinke et al. 2024](#)). Recent monitoring work has begun documenting analogous behaviors in deployed systems. The the Centre for Long-Term Resilience’s Loss of Control Observatory project analyzed more than 180,000 real-world user interactions on publicly accessible social media platforms and identified 698 instances displaying analogous scheming- and loss-of-control-related behaviors, and recorded a 4.9x increase in user-reported loss-of-control incidents over a five-month observation window ([Shane et al. 2026](#)). Behaviors that were once observable only in controlled environments are now appearing in real-world deployment, while the institutional capacity to surface and respond to them is still developing.

Child safety. Single deployment incidents producing harmful sexual abuse imageries have triggered investigations in multiple countries. These investigations have proceeded independently and produced different accounts and different outcomes for the same underlying incident ([Goggin 2025](#)). The documented failure modes include the generation of synthetic child sexual abuse material that spreads across borders through global platforms, and interaction-design choices that have amplified mental health crises among minors in every market the systems reach ([Thiel et al. 2023](#)). As a result, coordination to address these harms has been challenging because the investigative authority sits distributed across jurisdictions, and the connections between those investigations are still being built.

As harm rises across these and other risk categories, the governance challenge is one of geography as well as scale ([Bengio et al. 2026](#); [DSIT 2025](#)). The systems producing these harms

are developed in one jurisdiction, deployed in dozens of others, and governed by institutions whose authority stops at national borders. The next section examines the structural features of AI development and deployment that produce this jurisdictional mismatch and explains why addressing these harms at a domestic level is an incomplete solution.

2.2. The Systems Producing These Harms Span Jurisdictions

The AI systems producing AI harms, some of which we discuss in **Section 2.1**, are developed in a small number of jurisdictions and deployed across most of the world. In practice, the same incident often can involve multiple national governments at the same time. After exclusion from Project Glasswing, the Indian government held bilateral negotiations with both the United States and Anthropic to obtain equitable access to the Mythos model, citing the need to test its power grids, telecom networks, and banking systems against the class of AI-enabled cyber capabilities that the operation had exposed. Where a system built in one jurisdiction creates risks in another, the affected jurisdiction's main paths to act currently run primarily through diplomatic and commercial engagement with the developer and the developer's home government.

Three features of the AI supply chain produce this result.

Deployment geography. A GPAI model released in a developer jurisdiction reaches users in dozens of other countries within days, through application programming interfaces designed and governed in the developer's home market. A regulator examining a harmful output in any of those jurisdictions encounters a system whose training, fine-tuning, deployment decisions, and safety controls were determined elsewhere, under legal authority sitting in another government. The GTG-1002 operation, which targeted entities across multiple jurisdictions simultaneously, illustrates the scale at which AI models now operate. Response in that case was coordinated bilaterally through the developer, since the regulator-to-regulator channels that would otherwise carry it are still being built ([Anthropic 2025b](#)).

Supply chain geography. Training data is curated in one jurisdiction, compiled into datasets hosted in another, and used to train models in a third. Compute infrastructure, model weights, and inference capacity cross borders as a matter of architectural routine. This distribution places different stages of an AI system under different governments' authority, so when a failure surfaces at one stage, harm spreads through the AI supply chain across different sectors and countries, making remediation impossible. Liability regimes that assign distinct obligations to developers and deployers, such as the EU AI Act, address part of this problem within their own jurisdictions, while the cross-jurisdictional distribution of harm calls for arrangements that operate between regimes.

Asymmetric national capacity. Frontier AI development is concentrated in a handful of jurisdictions—primarily the United States and China—while these systems are deployed across most of the world. Most jurisdictions, and more disproportionately in the Global South, lack both the formal channels to access developer-held information and the domestic technical capacity to

analyze it if disclosed ([Wiaterek, Abungu and Okolo 2025](#)). The evidence base for CBRN risk, behavioral failure modes, and deployment-time anomalies is produced in ecosystems these jurisdictions cannot access. Cross-border infrastructure must extend deliberately to these jurisdictions, or it risks reproducing the asymmetries it is intended to address.

Together, these three features produce a cross-jurisdictional problem. When a powerful AI system causes harm, it is often unclear which jurisdiction's law applies, which law stands the best chance of working as intended, which authority is competent to investigate, which body can compel remediation, and which body can prevent the next incident. Each answer depends on the specific distribution of development, deployment, data, and harm across borders. Because that distribution differs for almost every incident, no single domestic regime, however well-designed, can resolve all five issues consistently.

Harmonizing domestic statutes is part of the answer but not the full answer. Shared thresholds and compatible evidentiary formats reduce the interoperability burden on developers and enable cross-jurisdictional comparison. What harmonization cannot supply on its own is the dialogue and protocol mechanisms through which disclosure in one jurisdiction can be recognized in another, investigative outputs can be mutually accepted, and verification of developer claims can be conducted, without duplicating effort across every jurisdiction with an interest in the outcome. That infrastructure must sit between jurisdictions.

2.3. Detection and Response for AI incidents Are Currently Reactive by Design

The capabilities of frontier AI systems are advancing on a timeline that the institutions tasked with detecting and responding to AI incidents have not been able to match. Present-day monitoring of AI incidents is constructed from voluntary disclosure, journalistic reporting, and public outcry—all of which are mechanisms that are reactive and often incapable of detecting patterns before the next serious incident materializes. The clearest expression of a push aimed at addressing this problem in existing incident reporting practices is the G7 Hiroshima AI Process Reporting Framework. This framework requests frontier AI developers to clearly disclose how they manage risk, to report AI incidents, and to share information ([OECD 2025a](#)).

At present, frontier AI companies are not obliged to publicly publish incident reports. Instead, all existing incident disclosure mechanisms are in practice voluntary in nature. However, voluntary AI incident disclosure is insufficient for anticipatory governance, since voluntary-only reporting omits incidents and produces inconsistent data that are not adequate for statistical analysis or accurately depicting real-world AI harms ([Dixon & Frase 2024](#); [Dixon & Frase 2025](#)). In the absence of consistent, legally mandated disclosure from frontier AI companies themselves, much of the work of documenting AI incidents has fallen to third parties with limited access to information. A growing body of technical research now tracks AI incidents, builds taxonomies of AI harms, and maps the landscape of AI risks ([McGregor 2021](#); [Hoffmann & Frase 2023](#); [Slattery](#)

[et al. 2024](#)). There are numerous public AI incident databases, including the [OECD AI Incidents and Hazards Monitor](#), the [AI Risk Explorer](#), the [AI Incident Database](#), and the [MIT AI Incident Tracker](#).

These databases, maintained by third-party volunteers, are fundamentally reliant on publicly available information, in that they compile and document incidents almost purely from public records and news sources. Such volunteer-led incident monitoring approaches are inadequate for comprehensive tracking and analysis of serious AI incidents. As a result, in addition to these efforts, some governments, such as [India](#) and [Japan](#), are also setting out national-level AI incident monitoring ambitions in their AI strategies and governance frameworks.

The aforementioned databases, taxonomies, datasets, and risk repositories that the field now relies on to catalogue incidents have largely been built by academics and think tanks working from public sources. More recently, governments, such as the European Union and California (United States), began pushing for mandatory reporting. These volunteer initiatives to track incidents are already showing results, surfacing a rising tide of harms, ranging from chatbot-linked suicides to autonomous agents hijacking cloud infrastructure ([OECD 2026](#)), and we are yet to see how much we will advance with the mandatory reporting in specific jurisdictions. The main challenge to having global data on AI harms, however, is that these initiatives are still nascent and jurisdiction-bound. Industry sits on the most granular data, including internal detection and early warning signals ([Kraprayoon 2024](#)) that in most cases never reach the public record. Accessing the relevant data has proven difficult, since companies are concerned with competitive sensitivity, legal exposure, user privacy, and the absence of trusted intermediaries.

As a result, visibility into AI harms today depends on voluntary disclosures from companies, public outcry around harms, and journalists reporting incidents that they can observe or access. There are a few emerging but currently immature efforts from certain actors to centralize the reporting of critical AI safety incidents and assessing threats to infrastructure. For example, following the legislative directives [SB 53](#) and [SB 896](#), the California Governor's Office of Emergency Services (Cal OES) must receive reports on AI-related incidents posing risks of death or serious injury within 24 hours. Similarly, the EU AI Office, via the [EU AI Act's Article 55\(1\)\(c\)](#) and the [GPAI Code of Practice's Commitment 9](#), has been explicitly designated as the official EU agency that frontier AI developers need to report serious incidents to for aggregate analysis, requiring disclosure within 2 days for widespread infringements or serious disruption of critical infrastructure, within 10 days for incidents resulting in death, and within 15 days for all other serious incidents. However, beyond such fragmented efforts, these data points have not been sufficiently systematized into trends nor addressed at an aggregate level. Furthermore, these mechanisms are reactive by design, and the systems producing the harm are advancing on a timeline they cannot match.

As a concrete example, in a report published in November 2025 ([Anthropic 2025](#)), Anthropic disclosed that they had detected ongoing use of their Claude Code model by a Chinese state-sponsored group. With the help of the AI tool, the group planned and orchestrated a highly

sophisticated cyber espionage campaign, where the attackers used AI not just as an advisor, but to execute the cyberattacks themselves. These attacks targeted major technology corporations, financial institutions, chemical manufacturing companies, and government agencies across multiple countries. Upon detection, the company was able to investigate these uses, respond to this malicious use of their services, and disrupt the attacks.

This is a successful example of effective internal incident infrastructure at work, where a company detected and responded to severe incidents within their own product ecosystem. This successful case, however, is limited by a few factors, including the fact that the disclosure was voluntary and that it was restricted to a single company and its model. A mandatory reporting system would not leave the provision of this information up to the company's will, and it would also enable responses across multiple services and ecosystems, promoting compounded benefits across multiple stakeholders and countries.

Notably, Anthropic's disclosure of the GTG-1002 operation, as we discussed above, has itself generated downstream demand for access to frontier cyber capabilities. India, for instance, initiated bilateral negotiations with the United States and Anthropic for equitable access to Anthropic's Mythos model, after India's exclusion from Project Glasswing. The Indian government cited the need to test the resilience of critical infrastructure, including power grids, telecom networks and banking systems, against the very class of AI-enabled vulnerabilities the GTG-1002 case exposed ([Agarwal 2026](#); [Anthropic 2026b](#)).

Cross-border AI incident infrastructure is the mechanism through which early signs can be surfaced, aggregated, and acted upon before the next serious incident materializes. In the next section, we will discuss in detail this infrastructure and provide recommendations for establishing it.

3. The Architecture Required and How It Can Be Secured

Robust cross-border AI incident infrastructure cannot be secured by any single institution or jurisdiction. It requires a set of interlocking mechanisms that together enable detection, analysis, sharing, and coordinated response across borders. The six components below constitute the necessary architecture, *each paired with specific actionable recommendations that decision-makers within national governments, AI Safety Institutes (AISIs), or multilateral institutions can take up within their existing institutional authority to help secure these architectural components.*

The incident infrastructure architecture we describe in this section does not refer to a single institution, but a series of mechanisms, spanning incident detection, prevention, preparedness,

and response. We have previously described such mechanisms by outlining solutions which are immediately actionable to existing institutions.² If secured, this architecture would facilitate greater incident resilience at a national level, and thereby consequently at the international level too, since the existence of effective national-level incident infrastructure is a necessary precondition for effective cross-border incident resilience.

1. Internationally interoperable protocols for incident documentation and reporting

An internationally shared reporting architecture is a necessary foundation on which robust cross-border AI incident infrastructure should be built. Absent internationally shared protocols for interoperable incident monitoring and reporting that enable reporting entities to clearly communicate the severity of an incident, hazard or near-miss, what categories of harm an incident falls under, and what informational fields a report should contain, incident data collected in one jurisdiction cannot be accessed, aggregated, compared, or analyzed alongside data collected in another. Each domestic regime ends up generating insights that are useful only within its own borders.

Since all jurisdictions require the same information from reported AI incidents, near-misses, hazards, flaws, and vulnerabilities for effective governance, an internationally shared incident reporting taxonomy that includes a standardized set of disclosed information would benefit all jurisdictions collectively. The interoperability produced by such a shared taxonomy is a precondition for effective cross-border incident pattern analysis, because it would minimize data gaps and maximize data quality. Moreover, it would create a shared language for understanding incidents and hazards, and a shared basis for standardized data structure, informational granularity, and severity classification across domestic reporting regimes, and would thereby make incident data comparable, aggregable, and analyzable across jurisdictions. This would enable researchers in academia, government, industry, and civil society to develop a comprehensive picture of the global AI incidents and hazards landscape, which they can use to more effectively prevent, prepare for, and respond to future incidents.

However, establishing such a foundation requires acknowledging a fundamental challenge: core terms like “AI incident” and “AI-enabled harm” are not yet universally defined ([Hoffmann & Frase 2023](#)). For example, being subjected to a privacy-violating AI system in the workplace may not be legally considered an “incident” or “harm” in jurisdictions focused strictly on provable physical or financial damages. Such sociotechnical disagreements make agreeing on a single, static incident taxonomy incredibly difficult. Because achieving perfect taxonomic consensus could take years, infrastructure and information-sharing mechanisms must be designed to operate even when the concept of an AI incident or harm is not perfectly shared. The solution lies in building systems that can communicate across differing definitions through shared technical architecture.

² For instance, the recommendations on serious accident prevention from [Herrmann & Lorente 2025](#); and the recommendations from the 2025 Athens Roundtable on AI and the Rule of Law ([The Future Society 2026](#)), and others.

Recommendation 1: *Governments and industry actors should adopt shared, machine-readable documentation structures and protocols to enable internationally interoperable systems for documenting and reporting information on AI incidents, near-misses, hazards, flaws, and vulnerabilities.*

Incident documentation structures and protocols that are universally compatible with one another are necessary to ensure that incident data collected in one jurisdiction becomes evidence that other jurisdictions can use. To achieve this in practice, jurisdictions could universally adopt the OECD's Common Reporting Framework (CRF) for AI Incidents ([OECD 2025b](#)). This framework was developed, vetted, and validated by an international group of experts and is ready to use. It provides the concrete informational fields that all incident reports should ideally aim to address to maximize informational completeness (e.g. what AI model was involved, who the model provider was, the number of affected parties) and gives both governments and industry a credible starting point to analyze incidents and cross-jurisdictional patterns.

Establishing a shared machine-readable structure to support interoperability and information sharing is also essential. For a robust cross-border incident infrastructure, systems across jurisdictions and organizations need to be able to “talk to” one another beyond basic taxonomic alignment. Current incident reporting forms typically produce non-standardized outputs that trap insights within domestic borders. Moving away from static, siloed databases, machine-readable incident report structures and output formats would enable (i) rapid AI processing of incident information, (ii) automated incident analysis and root cause identification, (iii) real-time information sharing during unfolding crises, and (iv) agile, internationally interoperable AI incident databases.

To achieve functional interoperability, differing jurisdictions can utilize policy baselines like the OECD CRF as a conceptual starting point, but must transition these into rigorous technical specifications. This work is already underway through technical standards development organizations; for instance, the European Telecommunication Standards Institute (ETSI) is developing technical specifications for compartmentalizing and sharing AI incident information ([ETSI TS 104 158-1/ETSI TS 104 158-2](#)), and the International Standards Organization (ISO) is advancing similar incident reporting standardizations ([ISO/IEC AWI 25870](#)). Moreover, in the US, the National Institute of Standards and Technology's (NIST) Center for AI Standards and Innovation (CAISI) is directed to “ensure AI is included in the establishment of standards, response frameworks, best-practices, and technical capabilities (e.g., fly-away kits) of incident response teams” ([US AI Action Plan. 2025](#)). Accordingly, it is currently pursuing initiatives such as convening a [multistakeholder workshop on AI incident management](#) to initiate a shared, multistakeholder dialogue on tractable approaches to AI incident response standards and practices, in an effort to shape a coordinated, forward-looking approach to AI incident management and implement America's AI Action Plan, including updates to existing guidelines, the development of new ones, and the strengthening of national and global alignment on AI incident preparedness and response. Even Commitment 9 (Measure 9.2) of the [EU's GPAI Code of Practice](#) appears to have taken inspiration from the OECD CRF to offer GPAI providers a basic

incident reporting taxonomy. These initiatives demonstrate how different jurisdictions can develop their own localized approaches that remain technically compliant and interoperable with global frameworks.

The ultimate value of adopting common incident reporting standards and these machine-readable, interoperable protocols is that they can enable a much broader ecosystem of information sharing and thereby enable aggregation across law enforcement agencies, consumer complaints, and industry reports to identify domestic and international patterns in AI-enabled harms. Currently, incident reporting is highly siloed, forcing practitioners to navigate a fragmented landscape of disjointed public databases, regulatory bodies, and industry consortia. By establishing internationally-shared reporting protocols, the ecosystem could support independent, interoperable incident reporting interfaces that allow a reporter to create a single incident report and seamlessly route it simultaneously to multiple incident databases (so that they all mirror the same incident record supported by the same incident information) and relevant stakeholders. To execute this safely and effectively, these interoperable systems should be explicitly designed to acknowledge and support the four distinct types of incident data: public and non-anonymized, public and anonymized, private and non-anonymized, and private and anonymized ([Brundage, 2020](#)).

Robust incident infrastructure requires specialized mechanisms tailored for both public and non-public incident data. Public-facing incident databases are vital for enabling accountability, authority notification, and global transparency. As a result, they should offer highly accessible reporting interfaces so that anyone can easily log incidents, near-misses, hazards, flaws, and vulnerabilities. In addition to the public incident databases listed in the previous section (e.g. the OECD's AIM, AIID), one leading operational example is [MITRE's AI Incident Sharing Initiative](#), which offers a voluntary, submission-based platform where any AI developer, deployer, researcher, or concerned citizen can anonymously submit reports on real-world AI incidents, failures, hazards, and near-misses. To encourage corporations to share their failures without fear of PR backlash while ensuring that incident data provides maximum value to policymakers and industry, MITRE acts as a neutral clearinghouse: it receives, sanitizes, and anonymizes incident data before distributing it only to a vetted trusted community of industry practitioners. It thus functions similarly to [Information Sharing and Analysis Centers \(ISACs\) from cybersecurity](#), one that welcomes public inputs and allows individuals to anonymously report incidents, near-misses, hazards, and vulnerabilities without public exposure while giving AI developers, deployers, and government agencies the situational awareness needed to spot emerging threats early. Ultimately, this initiative is vital because it transitions AI incident reporting from empirical incident observation to actionable operational security (practitioners sharing and acting upon granular, anonymized threat intelligence).

Equally critical are non-public, secure channels for sharing private data, which facilitate strategic industry coordination and rapid crisis response. A prime example is the Frontier Model Forum (FMF), which functions similarly to an [ISAC exclusively for frontier AI developers](#) by enabling the

[secure exchange of highly sensitive incident intelligence](#) that contains information hazards too dangerous for public release, such as incidents, hazards, near-misses, flaws, and vulnerabilities relating to CBRN capabilities, active jailbreaks, and cyber-threat indicators. The FMF offers the gold standard for secure intelligence exchange among frontier AI companies. To optimize the anticipatory governance value that could be derived from non-public incident sharing initiatives like this, policymakers, sectoral regulators, and industry leaders should establish additional ISAC-like intermediaries across diverse AI sectors, as proposed in the US's AI Action Plan ([White House 2025](#)). To complement such efforts, AI developers, deployers and researchers could start utilizing instruments like the [MIT-Stanford standardized AI flaw reporting framework](#) ([Longpre et al. 2025](#)), which aims to enable industry actors to collaboratively identify and address systemic AI flaws and vulnerabilities that affect multiple organizations across jurisdictions.

2. Mandatory incident documentation, disclosure, and preparedness obligations for frontier AI companies

Mandatory incident documentation, disclosure, and preparedness obligations for frontier AI companies are necessary to ensure that these practices are not contingent solely on the voluntary incident preparedness protocols and information-sharing cooperation of AI companies. It is essential to establish formal incident intelligence-sharing pipelines between AI companies and relevant government authorities. This measure would enable relevant national and international AI governance bodies to become aware of near-misses, hazards, and vulnerabilities. The expected result is the global governance system would be enabled to respond in a more timely manner to severe cross-sectoral and cross-jurisdictional AI harms that are currently only known by AI developers, thereby helping stakeholders understand risks from emerging capabilities much faster than their current capacities allow. Mandating AI companies to publish formal AI incident preparedness protocols would force them to establish rigorous procedures dedicated to ensuring the prevention of and effective management of AI risks, and could enable national governments to better understand and better prepare for future AI incidents that may occur within their borders.

Recommendation 2: *Jurisdictions that are home to frontier AI companies or with substantial AI deployment should enact statutory obligations for AI companies to report serious incidents to competent government authorities and to publish Serious Incident Prevention Policies.*

Relevant internal data held by frontier AI companies about serious AI incidents typically cannot be kept track of, gathered, and reported in retrospect after a serious incident has occurred, since information about the direct or indirect cause of such an event is often dispersed and may be lost, overwritten, or fragmented by the time AI developers become aware of a serious incident. As argued for in the EU's [GPAI Code of Practice Safety & Security Chapter](#), this justifies the establishment of processes and measures designed to keep track of and document relevant information before serious incidents occur.

For relevant internal incident information to reach competent government authorities and enable them to make effective governance decisions, AI incidents in each jurisdiction should be reported to a competent government agency or entity that has the capacity for meaningful incident investigation, to promote transparency and accountability in AI incident management. Ideally, this entity should be functionally independent of regulatory authorities, and should have the primary function of objectively investigating serious incidents and recommending preventive actions to frontier AI companies responsible for particular incidents and to relevant regulatory authorities. Moreover, this entity should commit to sharing relevant incident information, as well as resolutions whenever possible, with the public after a certain number of days or weeks has passed, to enable an informed public and consumer base.

Incident disclosure mechanisms should include the establishment of confidential reporting channels: dedicated, secure whistleblowing mechanisms that enable employees of frontier AI companies to confidentially report safety concerns relating to AI systems and models to relevant regulatory authorities. To prevent AI companies from opposing such disclosure obligations on privacy grounds and from fearing that government cooperation may lead to information leakage, the government authorities who would receive incident disclosures would need to commit to recognizing that the reporting of a serious incident is not an admission of wrongdoing, and to handling shared incident and hazard information with the same care they extend to national security-related information. Moreover, given that sharing incident data involves sharing affected users' private information, protecting the privacy of user data must also be prioritized. This could be done through the redacted or anonymized aggregation of user data prior to disclosure.

In its most robust form, incident infrastructure would ideally entail the adoption of a hybrid reporting framework. This framework should receive three types of reports. First, mandatory reporting, where relevant AI actors are mandated to report covered incidents in a timely manner. Second, it should also receive voluntary reporting, which comes from the good faith-based disclosure mechanisms, to capture additional AI incidents that fall outside mandatory reporting obligations. Finally, it should also be open to citizen reporting: accessible self-reporting mechanisms that enable the public and all other stakeholders to report and document AI incidents they observe or experience. This tripartite function would maximize data fidelity and provide an accurate and holistic representation of the emerging trends in AI harm and risk ([Dixon and Frase 2024](#)). This recommendation would address the first of these elements, which the stakeholders this report is aimed at can most readily act on.

Turning mandatory incident reporting obligations into reality does not require starting from scratch. China (through Article 14 of its [Interim Measures for Generative AI Services](#) and its more recent [National Cybersecurity Incident Reporting Management Measures](#)), the EU (through the EU AI Act's [Article 73](#)), Brazil (through [Bill No. 2338/2023](#)), Canada (through [Bill C-27](#)), and the United States (through California's [SB 53](#) and New York's [RAISE Act](#)) have already enacted or proposed guidelines for mandatory AI incident reporting obligations. Each of these existing

guidelines could be used as inspiration or starting points for other jurisdictions to enact or propose similar initiatives.

In addition to mandatory incident reporting obligations, governments should also require frontier AI companies to publish Serious Incident Prevention Policies (SIPP, [The Future Society 2025](#)) that demonstrate their preparedness for preventing and effectively responding to serious AI incidents. These are documented policies by frontier AI companies setting out specific personnel roles and responsibilities, incident monitoring mechanisms, emergency shutdown protocols, and pre-established plans, outlining the measures to be taken in the event of serious AI incidents to contain and control their negative effects. Multilateral institutions could play a part in encouraging the national enactment of SIPP across member and partner countries.

Modeled on major accident prevention frameworks from other high-stakes industries—such as [the EU’s Seveso-III Directive](#), the UK’s [COMAH15](#) regulations, and the US’s [OSHA Process Safety Management 29 CFR 1910.119](#) and [EPA Risk Management 40 CFR Part 68](#) programs—SIPP would require frontier AI companies to document, disclose, test, and periodically update their protocols for preventing, containing, and recovering from serious AI incidents and near-misses. This includes implementing protective measures, stress-testing potential incident scenarios and response plans (e.g. modeling scenarios and responses via table-top exercises), and communicating necessary information to the public and relevant governmental and regulatory authorities. This architectural component would be made even more robust if governments establish independent SIPP oversight bodies: independent agencies with the responsibility of overseeing and enforcing the implementation of AI companies’ serious incident prevention policies. If they do so, these oversight bodies could share relevant information with one another across jurisdictions, which could enable national governments to better understand, and better prepare for, future AI incidents that may occur within any of their borders.

3. Domestic incident data access and institutional coordination

For incident-relevant data to usefully inform policy action, national AISIs and functionally equivalent institutions need to have access to information from all data sources that are relevant to incident prevention and preparedness, which goes beyond the internal incident data that frontier AI companies can provide. This would ensure that their work is informed by an evidence base recorded by the widest possible range of actors involved in the AI development and deployment lifecycle.

Recommendation 3: *Governments should grant their national AISI-equivalents access to domestic incident data from non-frontier-AI-company actors within their jurisdictions and establish formal information-sharing arrangements between AISI-equivalents and appropriate regulatory authorities.*

It is crucial for effective AI governance, both national and international, that government support their national AISIs and functionally equivalent counterparts in securing access to relevant incident data held by AI deployers and sectoral regulators (for instance, financial services, healthcare, critical infrastructure, consumer protection, online safety) operating within their jurisdictions. Many AI deployers collect and store crucial internal safety telemetry concerning their internal and external AI deployments. Similarly, sectoral regulators typically have access to a wide array of incident data—ranging from immediate operational safety or security incident reports (e.g. that detail technology failures, cyberattacks, or service disruptions such as network outages) to aggregated data on consumer complaints and feedback—gathered through a variety of channels such as direct reporting (e.g. via dedicated digital portals), mandatory data submissions (i.e. via regular, scheduled reports like quarterly data on consumer complaints), and third party sources (i.e. via information from media, social media, or other government departments).

Ensuring that these institutions have the statutory authority to access important incident data held by these actors, subject to confidentiality protections, would enable them to convert their existing technical expertise into usable analytical output. To ensure this is achieved in practice, governments should establish (or support the establishment of) voluntary national AI incident reporting hubs and databases aimed at aggregating cross-sector data and detecting distributed harms before they escalate. In the US, multiple organizations ([Center for Data Innovation](#), [FAS](#), [ARI](#), [CNAS](#)) already support the establishment of a voluntary national AI incident database modeled on the FDA/FAA systems, where NIST could host confidential data collection and allow AI deployers to report incidents without public exposure while giving government agencies the situational awareness needed to spot emerging threats early ([Jeanmaire and Boger 2025](#)). Once a certain number of days or weeks has passed following an incident, these government institutions could then share relevant findings with the public to inform them of materialized AI harms.

To be effectively realized, this architectural component of incident infrastructure would also involve the establishment of formal, national-level information-sharing arrangements between these institutions and the relevant regulatory authorities who can act upon the insights these institutions can provide. These arrangements could ensure that relevant incident data reaches actors who can take specific regulatory actions in response to them within the jurisdictions in which incidents occur, enabling these bodies to devise more up-to-date and future-proof governance mechanisms. Just as with the incident disclosure channels proposed in the previous recommendation, these arrangements would also need to ensure adequate privacy protections for both the AI deployers and regulators who share incident-relevant data and for the victims affected by reported incidents.

The main drivers behind this recommendation are that (i) meaningful monitoring requires independent access to non-public technical and organizational information ([Brundage et al. 2020](#)), and (ii) where developers make such information available, they typically only do so under

domestic regulatory obligation or by bespoke arrangements. Standards for how frontier AI developers, downstream deployers, sectoral regulators, AISIs, and government agencies should work with and share safety and security-relevant information with one another in a timely manner, and what this information should include, are still nascent. This architectural component and linked recommendation would help address this critical gap.

It would make sense for AISI-equivalent institutions to be the ones to be granted such information access because they possess the technical capacities needed to meaningfully analyze complex incidents. Many of them have privileged early access to frontier AI systems (for testing), deep expertise, and existing legal arrangements that can protect sensitive industry information without jeopardizing their own independence or companies' competitiveness, and so they are well-positioned to derive useful insights from granular incident data.

4. Adequate international incident analysis capacity and expertise

Sufficient incident analysis capacity and expertise is a necessary condition for maximally useful and actionable insights to be extracted from documented AI incidents, and for these insights to be used by various stakeholder groups, both domestically and internationally, to prevent and prepare for future incidents.

Recommendation 4: *International standards development organizations should develop a standardized incident root-cause and supply-chain analysis methodology that can be internationally adopted, actively advocated for, and promoted by multilateral institutions.*

One of the primary constraints limiting effective international incident response is the absence of a common analysis methodology. Accordingly, much like the OECD's aforementioned CRF, which standardizes the reporting of data, this proposed initiative would focus on standardizing the analysis of that data. The international adoption of such a methodology would help deepen incident analysis capacities and expertise both nationally and internationally.

In an ideal world, an independent agency would be tasked with AI incident investigation and analysis. Their role would involve investigating the root causes of AI incidents documented around the world and providing evidence-based safety recommendations to relevant national government authorities. This would enable severe cross-border incidents to be analyzed in a centralized manner. The reason why relying solely on ad-hoc cooperation between national AI governance agencies would be less desirable than the establishment of such an independent agency is because national agencies inherently face domestic political pressures, industrial policy incentives to protect their own domestic AI champions, and conflicting geopolitical agendas. In the face of this reality, a standalone, independent agency would ensure objective, neutral, and politically uncompromised investigations, thereby fostering the international trust required for companies to confidently share sensitive operational data.

However, given the costs and time associated with creating such an institution, an alternative pathway would be for an international standards development organization (e.g. ISO, IEC, CEN-CENELEC) to develop a standardized AI incident root-cause and supply-chain analysis methodology that is advocated for and actively promoted by multilateral institutions (e.g., the OECD or UN). After all, standards development organizations already possess the technical expertise and robust multistakeholder processes required to develop the complex components that are likely to comprise the standardized root-cause and supply-chain analysis methodology in question. This work could build on existing foundational initiatives in this vein, such as the MIT AI Risk Initiative ([MIT 2025](#)), which provides stakeholders with a structured approach for analyzing and extracting policy-relevant insights from documented incidents. Throughout the development process, frontier AI companies would need to be encouraged to share their current incident analysis practices and translate their unique knowledge and incident data into transferable analytical methodologies.

The development and implementation of this methodology could be made significantly more practical through the universal adoption of a common reporting framework for AI incidents, as called for by Recommendation 1. If frontier AI developers, government agencies, academic researchers, and civil society organizations are all actively encouraged to document information about AI incidents, near-misses, hazards, and vulnerabilities using the same or similar incident reporting and documentation frameworks, when incidents inevitably occur and root-cause analyses are required, the informational foundations and common technical vocabulary required to analyze them would already be in place. This starting point would ensure that analytical insights can be produced much more quickly and understood more widely across different jurisdictions.

Typically, analytical methodologies are developed and adopted in a single jurisdiction or industry consortium before spreading globally. However, multilateral institutions like the OECD and UN could play a vital role in encouraging and speeding up the international adoption of this methodology. They could achieve this through several concrete mechanisms. For example, by actively advocating for technical AI governance as a recommended practice, they could prompt national policymakers to direct the energy and resources of their domestic standards bodies toward participating in the development of shared methodologies. Moreover, they could consistently prioritize this topic in all their multilateral discussion forums (e.g. UN Global Dialogue on AI Governance, G7 tracks, AI Safety Summits), formally establish and regularly convene dedicated expert working groups on incident analysis, and create confidential, safe-harbor technical sub-committees where industry scientists can candidly share post-mortem practices without fear of immediate regulatory or reputational reprisal. After all, these organizations possess the legitimacy and cross-jurisdictional reach needed to convene the diverse multi-stakeholder group of international experts required to devise a sufficiently robust methodology.

Once established, this methodology could be immediately taken up and broadly adopted by a diverse ecosystem of actors across their respective jurisdictions, spanning government authorities, frontier AI companies, academic researchers, and civil society organizations. Nationally competent AI governance agencies, such as AISIs and their functional equivalents, would likely serve as the methodology's primary implementers, as they already possess the technical expertise required to extract meaningful insights from complex incident data. Sectoral regulators would be similarly strong candidates for implementation, given their unique access to sector-wide data crucial for deep supply-chain and root-cause analysis. Regardless of what specific national bodies are assigned to adopt such a methodology, governments should invest heavily in expanding the technical capacities required to apply it in practice.

5. Cross-border incident-sharing channels and joint investigation mechanisms

Formalized cross-border incident-sharing channels and joint investigation mechanisms are necessary to ensure that incident types that tend to materialize across jurisdictions can be identified and acted upon in an internationally coordinated manner. This can enable governments and internationally distributed industry actors to track concerning AI incidents and risk categories across jurisdictions. Consequently, this allows us to understand similarities and differences between incidents found in different countries and contexts. Ultimately, it strengthens the collective incident prevention and preparedness capacities of AI governance bodies globally, helping them come up with more internationally coherent AI governance mechanisms.

Recommendation 5: *Governments should enable their national AISIs and functionally equivalent agencies to exchange incident intelligence with international counterparts.*

The international exchange of incident data for cross-border incidents that exceed the investigative capacity of any single jurisdiction would enable governments around the world to learn collectively from such incidents, rather than each trying to learn purely from the isolated harms they individually encounter within their jurisdictions. By establishing defined multi-state and multistakeholder incident-sharing procedures, governments and industry actors could support each other in building up their collective incident resilience capacities. Such international cooperation could, for example, be achieved through the widespread use of [MITRE's AI Incident Sharing Initiative](#) by industry actors and/or through the establishment of an international public-private ISAC that is dedicated exclusively to the sharing of information that relates to severe cross-border AI incidents, near-misses, and hazards.

Acknowledging that AI companies may not want incident information to be widely shared among international government agencies, and that government agencies may not want to share particularly sensitive national security-related information, certain kinds of information could be redacted prior to sharing.

The International Network for Advanced AI Measurement, Evaluation and Science could be one promising avenue through which such a cross-border incident information-sharing regime could begin. This Network could operationalize incident sharing among its member institutes, beginning with bilateral arrangements between willing members, which would move the architecture from aspiration to practice. Alternatively, a new international mechanism could be proposed through the UN's Global Dialogue on AI Governance.

The [ICAO Annex 13 framework](#) for joint investigation from the international civil aviation sector offers a directly applicable model for inspiration here, covering confidentiality rules, defined roles for the state of occurrence and the state of the operator, and mutual recognition of investigative outputs. A general agreement among all jurisdictions would not be required as a starting point for this to be achieved; progress could be made through bilateral instruments that others could later join as their domestic arrangements mature.

6. Internationally distributed incident detection, analysis, and response capacity

Robust cross-border incident infrastructure requires *resilient*, internationally distributed incident detection, analysis, and response capacities. If incidents are detected, analyzed, and responded to effectively in only one jurisdiction and not in other jurisdictions, AI systems will continue to cause harm.

Recommendation 6: *Multilateral institutions should help extend incident infrastructure to jurisdictions who have inadequate domestic incident monitoring and management capacities.*

Many of the communities most acutely exposed to AI harms are also those with least facing the least coverage in existing incident monitoring systems and with the weakest capacities for preventing, preparing for, and effectively responding to AI incidents. These include users interacting with models in low-resource languages, those in jurisdictions with weaker press freedoms, and those affected by enterprise deployments that never reach public incident databases. By encouraging and actively supporting nations and organizations in expanding AI incident infrastructure to more jurisdictions, multilateral institutions could help secure a maximally robust and comprehensive cross-border incident management regime.

Securing the inclusion of a greater range of jurisdictions in the development of cross-border AI incident infrastructure is a matter of mutual interest. Since most AI models are deployed internationally across a wide range of application domains, incident findings in all jurisdictions, like the discovery that GPT-4 would readily provide users with bomb-building advice when prompted in Zulu ([New Scientist 2023](#)), serve as an alarm for the whole AI governance system.

Practical mechanisms to achieve this could include the establishment of bilateral or, where appropriate, multilateral intelligence-sharing agreements between national AISI-equivalents; capacity building partnerships; or incident response coordination plans between different

countries or regional blocs. We are aware that there are many countries with weak AI governance regimes are unlikely to have AISI-equivalent government agencies. One solution could be for nearby countries who already possess relatively mature incident infrastructure to start sharing the incident information they process, as well as useful guidance on effectively using this information to prepare for and respond to future incidents. This would create collaboration between governments in the same region who could pool domestic capacities.

It is important to note that none of the six architectural components and recommendations described above are self-sufficient in isolation. While each can be acted upon independently and bilaterally, and will certainly advance AI governance in each of their relevant scopes and mechanisms, they are interlocking by design. Each addresses a specific gap identified in Section 2, and each is strengthened by the others being in place. If each of the above recommendations are effectively acted upon, the prevention and management of AI incidents would have the foundational infrastructure necessary to substantially improve, both domestically and internationally.

4. Conclusion

Today, the mechanisms for detecting and responding to AI incidents sit largely inside national borders or inside the companies that build the systems, forcing a difficult choice between waiting for each jurisdiction to build its own response capacity and relying on whatever portion of an incident a developer chooses to disclose through bilateral channels. Cross-border infrastructure provides an alternative that could close this gap through a coordinated set of institutional, technical, and legal arrangements built between national and international institutions.

Through this report, we presented a vision for what effective cross-border AI incident infrastructure requires. These include **interoperable protocols** that allow incident data collected in one jurisdiction to be used in another; **mandatory documentation, disclosure, and preparedness obligations** on the companies building frontier systems; **statutory access for national AISIs and functionally equivalent bodies** to incident data held by deployers and sectoral regulators within their jurisdictions; **a shared analytical methodology** that allows root causes and supply-chain dynamics to be studied across borders; **formal cross-border channels** through which AISIs and their counterparts can exchange incident intelligence and conduct joint investigations; and **deliberate extension of this architecture to all jurisdictions** currently outside the systems that monitor and analyze AI harm. These components draw on more established domains where governments have repeatedly built shared infrastructure for managing risks that no single jurisdiction could contain alone.

Real difficulties stand in the way. The field still lacks a settled definition of what counts as an AI incident, hazard, or crisis. Frontier developers hold information whose disclosure carries genuine commercial and security costs. The asymmetry between jurisdictions is wide enough that infrastructure built without care will deepen it. These are familiar shapes of problem, and much effort is needed to learn from other safety-critical domains that have worked through versions of each dilemma, to draw on what they learned.

Acknowledgments

The authors would like to thank **Abra Ganz, Kevin Path, Nick Moës, Niki Iliadis, Peter Slattery, Shayne Longpre, and Simon Mylius** for their invaluable feedback on earlier drafts of this report. We are grateful to **Benedicte Rispal, Luis Aranda, Sean McGregor, and Tommy Shaffer Shane**, who have enriched the development of these ideas. All mistakes remain our own.

We also wish to acknowledge the participants of the closed-door dialogue session on “[Serious AI Incident Prevention and Preparedness](#)” that we held at the 7th Edition of The Athens Roundtable, and participants of our workshop at the India AI Impact Summit, “[Closing the Loop: Building Global AI Incident Monitoring and Response Capabilities](#)”, for the rich discussions that informed key aspects of this work, as well as the broader AI incidents community for all their contributions to the field, which directly inspired this report.

References

- Agarwal, Surabhi. 2026. "India Seeks Fair Access to Anthropic's Mythos for Critical Infrastructure Security." *Economic Times*.
<https://economictimes.indiatimes.com/tech/artificial-intelligence/india-seeks-fair-access-to-anthropics-mythos-for-critical-infrastructure-security/articleshow/130588755.cms>.
- Anthropic. 2025a. "Protecting the Well-Being of Our Users."
<https://www.anthropic.com/news/protecting-well-being-of-users>.
- Anthropic. 2025b. "Disrupting the First Reported AI-Orchestrated Cyber Espionage Campaign." November 13. <https://www.anthropic.com/news/disrupting-AI-espionage>. Full report: <https://assets.anthropic.com/m/ec212e6566a0d47/original/Disrupting-the-first-reported-AI-orchestrated-cyber-espionage-campaign.pdf>.
- Anthropic. 2026a. "Risk Report: February 2026."
<https://www-cdn.anthropic.com/08eca2757081e850ed2ad490e5253e940240ca4f.pdf>.
- Anthropic. 2026b. "Project Glasswing: Securing Critical Software for the AI Era." April 7.
<https://www.anthropic.com/glasswing>.
- Artificial Intelligence Incident Database via AI Index Report. 2026. "Global annual number of reported artificial intelligence incidents and controversies." Our World in Data.
<https://ourworldindata.org/grapher/annual-reported-ai-incidents-controversies>.
- Bengio, Yoshua, et al. 2026. *International AI Safety Report 2026*. February.
https://internationalaisafetyreport.org/sites/default/files/2026-02/international-ai-safety-report-2026_1.pdf.
- Brundage, Miles, Shahar Avin, Jasmine Wang, Haydn Belfield, Gretchen Krueger, Gillian Hadfield, Heidy Khlaaf, et al. 2020. "Toward Trustworthy AI Development: Mechanisms for Supporting Verifiable Claims." arXiv:2004.07213. <https://arxiv.org/abs/2004.07213>.
- Brundage, Miles, Noemi Dreksler, Aidan Homewood, Sean McGregor, Patricia Paskov, Conrad Stosz, Girish Sastry, A. Feder Cooper, George Balston, Steven Adler, et al. 2026. "Frontier AI Auditing: Toward Rigorous Third-Party Assessment of Safety and Security Practices at Leading AI Companies." AVERI. <https://www.averi.org/ourwork/frontier-ai-auditing>.
- Cress, Laura. 2026. "Family of Child Injured in Canada School Shooting Sues OpenAI." *BBC News*, March 10. <https://www.bbc.co.uk/news/articles/c309y25prnlo>.
- Dixon, Ren Bin Lee, and Heather Frase. 2024. "An Argument for Hybrid AI Incident Reporting." Center for Security and Emerging Technology, March.
<https://cset.georgetown.edu/publication/an-argument-for-hybrid-ai-incident-reporting/>.
- DSIT (Department for Science, Innovation and Technology). 2023. *Capabilities and Risks from Frontier AI: A Discussion Paper on the Need for Further Research into AI Risk*. London: UK

- Government, October.
<https://www.gov.uk/government/publications/frontier-ai-capabilities-and-risks-discussion-paper>.
- Elagina, Diana. 2026. “Number of AI Incidents and Hazards Worldwide Monthly 2020–2026.” Statista. <https://www.statista.com/statistics/1659751/ai-incidents-and-hazards-worldwide/>.
- French Republic. 2025. “Sommet pour l’action sur l’IA / AI Action Summit.” February 10–11. Élysée. <https://www.elysee.fr/en/sommet-pour-l-action-sur-l-ia>.
- Gardhouse, Kathrin. 2026. “An Opportunity for Canada to Lead in AI Emergency Preparedness.” The Future Society. <https://thefuturesociety.org/canada-ai-emergency-preparedness>.
- Girrus, Peter, Vincenzo Ciancaglini, Morton Swimmer, David Fiser, Alfredo Oliveira, Benjamin Zigh. 2026. “Fault Lines in the AI Ecosystem: TrendAI™ State of AI Security Report.” Trend Micro.
<https://www.trendmicro.com/vinfo/gb/security/news/threat-landscape/fault-lines-in-the-ai-ecosystem-trendai-state-of-ai-security-report>.
- Goggin, Ben. 2025. “Possession of AI-Generated Child Sexual Abuse Imagery May Be Protected by First Amendment in Some Cases, Judge Rules.” *NBC News*.
<https://www.nbcnews.com/tech/tech-news/ai-generated-child-sexual-abuse-imagery-judge-ruling-rcna196710>.
- Gor, George, and Niki Iliadis. 2025. “What Is an Artificial Intelligence Crisis and What Does It Mean to Prepare for One?” The Future Society.
<https://thefuturesociety.org/aicrisisexplainer/>.
- Government of India, Ministry of External Affairs. 2026. “India AI Impact Summit: Bilateral Documents.” <https://www.mea.gov.in/bilateral-documents.htm?dtl/40809>.
- Herrmann, Sven, and Toni Lorente. 2025. “Serious Incident Prevention for AI: Lessons from Other Industries and Recommendations for the EU AI Office.” The Future Society.
<https://thefuturesociety.org/serious-ai-incident-prevention-eu/>.
- Hoffmann, Mia, and Heather Frase. 2023. “Adding Structure to AI Harm: An Introduction to CSET’s AI Harm Framework.” Center for Security and Emerging Technology, July.
<https://cset.georgetown.edu/publication/adding-structure-to-ai-harm/>.
- Ileri, Gathoni. 2025. “What We Risk When AI Systems Remember.” *Tech Policy Press*, October 21.
<https://www.techpolicy.press/what-we-risk-when-ai-systems-remember/>.
- Ileri, Gathoni, Cecil Abungu, Jean Cheptumo, Sienka Dounia, Mark Gitau, Stephanie Kasaon, Michael Michie, Chinasa T. Okolo, et al. 2026. “Assessing the Case for Africa-Centric AI Safety Evaluations.” arXiv:2602.13757. <https://arxiv.org/abs/2602.13757>.

- Jeanmaire, Caroline, and Sam Boger. 2025. "AI Incidents Are Rising. It's Time for the United States to Build Playbooks for When AI Fails." The Future Society.
<https://thefuturesociety.org/us-ai-incident-response/>.
- Judd, Amy. 2026. "Tumbler Ridge B.C. School Shooting: 8 Dead, 27 Injured." *Global News*, February 11. <https://globalnews.ca/news/11662006/tumbler-ridge-school-shooting/>.
- Kraeva, Radina, Amin Oueslati, and Toni Lorente. 2026. "Europe's AI Strategy: Mapping the EU's Emerging AI Policy Portfolio." The Future Society.
<https://thefuturesociety.org/mapping-europes-emerging-ai-policy-strategy>.
- Kraprayoon, Jam, Joe O'Brien, and Shaun Ee. 2024. "An Early Warning System for AI-Powered Threats to National Security and Public Safety." Institute for AI Policy and Strategy, June.
<https://www.iaps.ai/research/early-warning-system>.
- Longpre, Shayne, Kevin Klyman, Ruth E. Appel, Sayash Kapoor, Rishi Bommasani, Michelle Sahar, Sean McGregor, et al. 2025. "In-House Evaluation Is Not Enough: Towards Robust Third-Party Flaw Disclosure for General-Purpose AI." arXiv:2503.16861.
<https://arxiv.org/abs/2503.16861>.
- Luckey, David, Sara Duhachek Muggy, Taylor Frey, David Stebbins, Tracey Rissman, Bianca Espinosa, Daniel Tapia, Greg McKelvey Jr., Neeti Pokhriyal, Joseph Nelson Dawson, et al. 2025. *Mitigating Risks at the Intersection of Artificial Intelligence and Chemical and Biological Weapons*. RR-A2990-1. Santa Monica, CA: RAND Corporation.
https://www.rand.org/pubs/research_reports/RRA2990-1.html.
- McGregor, Sean. 2021. "Preventing Repeated Real World AI Failures by Cataloging Incidents: The AI Incident Database." *Proceedings of the AAAI Conference on Artificial Intelligence* 35 (17): 15458–63. <https://doi.org/10.1609/aaai.v35i17.17817>.
- Meinke, Alexander, Bronson Schoen, Jérémy Scheurer, Mikita Balesni, Rusheb Shah, and Marius Hobbhahn. 2024. "Frontier Models Are Capable of In-Context Scheming." arXiv:2412.04984. <https://arxiv.org/abs/2412.04984>.
- MIT AI Risk Initiative. 2025. *AI Risk Repository*. Massachusetts Institute of Technology.
<https://airisk.mit.edu/>.
- New Scientist. 2023. "GPT-4 Gave Advice on Planning Terrorist Attacks When Asked in Zulu." October.
<https://www.newscientist.com/article/2398656-gpt-4-gave-advice-on-planning-terrorist-attacks-when-asked-in-zulu/>.
- NIST (National Institute of Standards and Technology). 2026. "International Network for Advanced AI Measurement, Evaluation, and Science Publishes Consensus Areas on Practices for Automated Evaluations." February.

<https://www.nist.gov/news-events/news/2026/02/international-network-advanced-ai-measurement-evaluation-and-science>.

OECD. 2024. “Defining AI Incidents and Related Terms.” OECD Artificial Intelligence Papers, No. 16. Paris: OECD Publishing. <https://doi.org/10.1787/d1a8d965-en>.

OECD. 2025a. “OECD Launches Global Framework to Monitor Application of G7 Hiroshima AI Code of Conduct.” Press release, February 7. <https://www.oecd.org/en/about/news/press-releases/2025/02/oecd-launches-global-framework-to-monitor-application-of-g7-hiroshima-ai-code-of-conduct.html>.

OECD. 2025b. *Towards a Common Reporting Framework for AI Incidents*. OECD Artificial Intelligence Papers, No. 34. Paris: OECD Publishing. https://www.oecd.org/en/publications/towards-a-common-reporting-framework-for-ai-incidents_f326d4ac-en.html.

OECD. 2026. *Trends in AI Incidents and Hazards Reported by the Media*. OECD Artificial Intelligence Papers, No. 53. Paris: OECD Publishing. <https://doi.org/10.1787/4f5ff43c-en>. See also the OECD AI Incidents and Hazards Monitor (AIM): <https://oecd.ai/en/incidents>.

OECD.AI. 2026. “Alibaba AI Agent ROME Engages in Unauthorized Crypto Mining and Network Tunneling.” OECD AI Incidents and Hazards Monitor, March 7. <https://oecd.ai/en/incidents/2026-03-07-95e2>.

OpenAI. 2025. “Our Approach to Mental Health-Related Litigation.” November 25. <https://openai.com/index/mental-health-litigation-approach/>.

Paeth, Kevin, Daniel Atherton, Nikiforos Pittaras, Heather Frase, and Sean McGregor. 2025. “Lessons for Editors of AI Incidents from the AI Incident Database.” *Proceedings of the AAAI Conference on Artificial Intelligence* 39 (28): 28946–53. <https://doi.org/10.1609/aaai.v39i28.35163>.

Republic of Korea. 2024. “Seoul Ministerial Statement for Advancing AI Safety, Innovation and Inclusivity.” May 22. <https://www.gov.uk/government/publications/seoul-ministerial-statement-for-advancing-ai-safety-innovation-and-inclusivity-ai-seoul-summit-2024/seoul-ministerial-statement-for-advancing-ai-safety-innovation-and-inclusivity-ai-seoul-summit-2024>.

Shaffer Shane, Tommy, Simon Mylius, and Hamish Hobbs. 2026. *The Loss of Control Observatory: A Prototype to Detect Real-World AI Control Incidents*. Centre for Long-Term Resilience. <https://www.longtermresilience.org/reports/the-loss-of-control-observatory-a-prototype-to-detect-real-world-ai-control-incidents/>.

- Sharma, Mrinank, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askill, Samuel R. Bowman, Newton Cheng, et al. 2023. "Towards Understanding Sycophancy in Language Models." arXiv:2310.13548. <https://arxiv.org/abs/2310.13548>.
- Slattery, Peter, Alexander K. Saeri, Emily A. C. Grundy, Jess Graham, Michael Noetel, Risto Uuk, James Dao, Soroush Pour, Stephen Casper, and Neil Thompson. 2024. "The AI Risk Repository: A Comprehensive Meta-Review, Database, and Taxonomy of Risks from Artificial Intelligence." arXiv:2408.12622. <https://arxiv.org/abs/2408.12622>.
- The Future Society. 2025. "Serious Incident Prevention for AI: Lessons from Other Industries and Recommendations for the EU AI Office." <https://thefuturesociety.org/serious-ai-incident-prevention-eu/>.
- The Future Society. 2026. *Serious AI Incident Prevention and Preparedness: Athens 2025 Roundtable Report*. February. <https://thefuturesociety.org/wp-content/uploads/2026/02/Serious-AI-Incident-Prevention-Preparedness-Athens-2025-Report.pdf>.
- Thiel, David, Melissa Stroebe, and Rebecca Portnoff. 2023. *Generative ML and CSAM: Implications and Mitigations*. Stanford Internet Observatory and Thorn, June 24. <https://purl.stanford.edu/jv206yg3793>.
- UK AI Security Institute. 2026. "Our Evaluation of Claude Mythos Preview's Cyber Capabilities." AISI Blog, April. <https://www.aisi.gov.uk/blog/our-evaluation-of-claude-mythos-previews-cyber-capabilities>.
- UK Government. 2023. "AI Safety Summit 2023." Topical event. <https://www.gov.uk/government/topical-events/ai-safety-summit-2023>.
- United Nations. 2025a. "Independent International Scientific Panel on AI." <https://www.un.org/independent-international-scientific-panel-ai/en>.
- United Nations. 2025b. "Global Dialogue on AI Governance." <https://www.un.org/global-dialogue-ai-governance/en>.
- White House. 2025. *Winning the Race: America's AI Action Plan*. Washington, DC: The White House, July 23. <https://www.whitehouse.gov/wp-content/uploads/2025/07/Americas-AI-Action-Plan.pdf>.
- Wiaterek, Joanna, Cecil Abungu, and Chinasa T. Okolo. 2025. "Building Regional Capacity for AI Safety and Security in Africa." Brookings. <https://www.brookings.edu/articles/building-regional-capacity-for-ai-safety-and-security-in-africa/>.
- Yousif, Nadine. 2026. "Canada Summons OpenAI Senior Staff over Tumbler Ridge Shooting." BBC News. <https://www.bbc.co.uk/news/articles/c2e4nvjwnno>.

THE
FUTURE
SOCIETY

Contact Us

GENERAL info@thefuturesociety.org

PRESS press@thefuturesociety.org



THE FUTURE SOCIETY

867 Boylston Street
Boston, MA, 02116
United States

www.thefuturesociety.org

